



The Evolution of the DevOps Pipeline: To K8s and Beyond!

Josh Lynn – Senior Cloud Engineer - LLamasoft



Josh Lynn

- 15 years in Industry
 - Thomson Reuters
 - Hooklogic (Acquired by Criteo)
 - Criteo
 - LLamasoft



Agenda

Introduction

By the Numbers

Packer, Chef, Terraform, Sequencer, Consul

Jenkins Pipelines

Kops

Kubernetes (K8s)

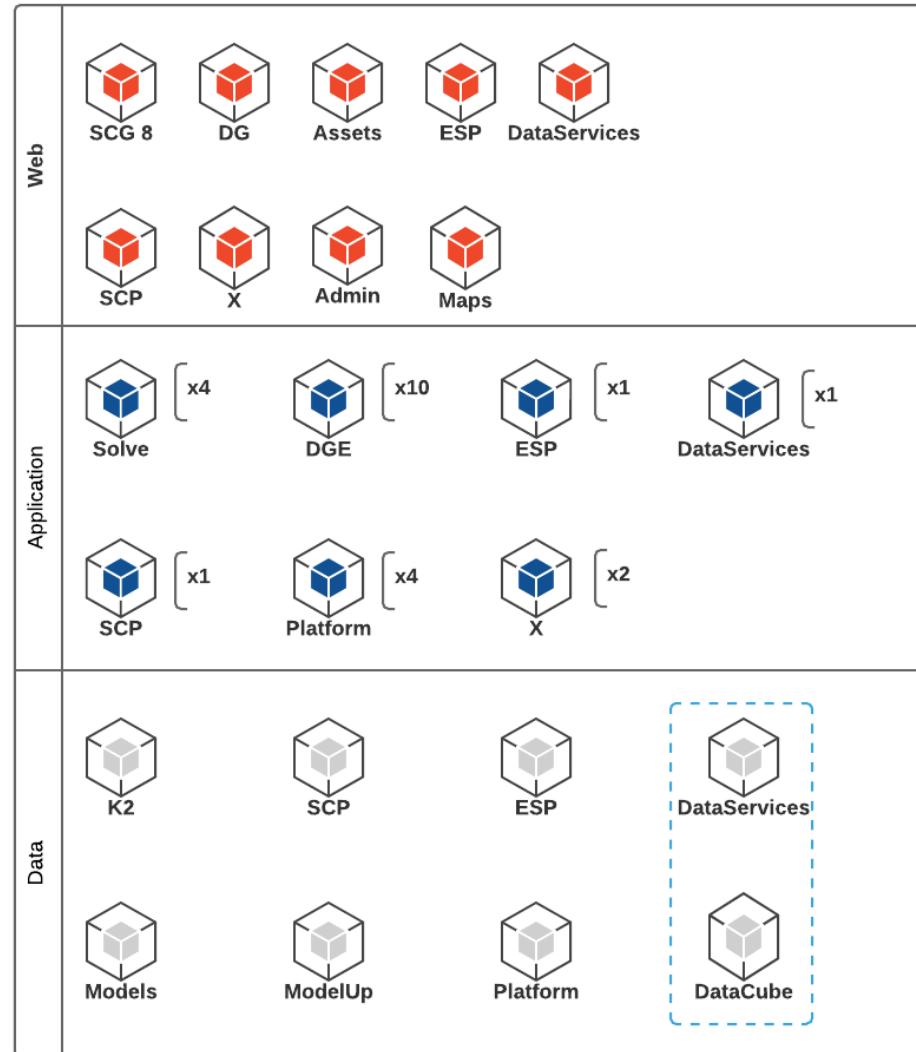
K8s Lessons Learned

What's Next?

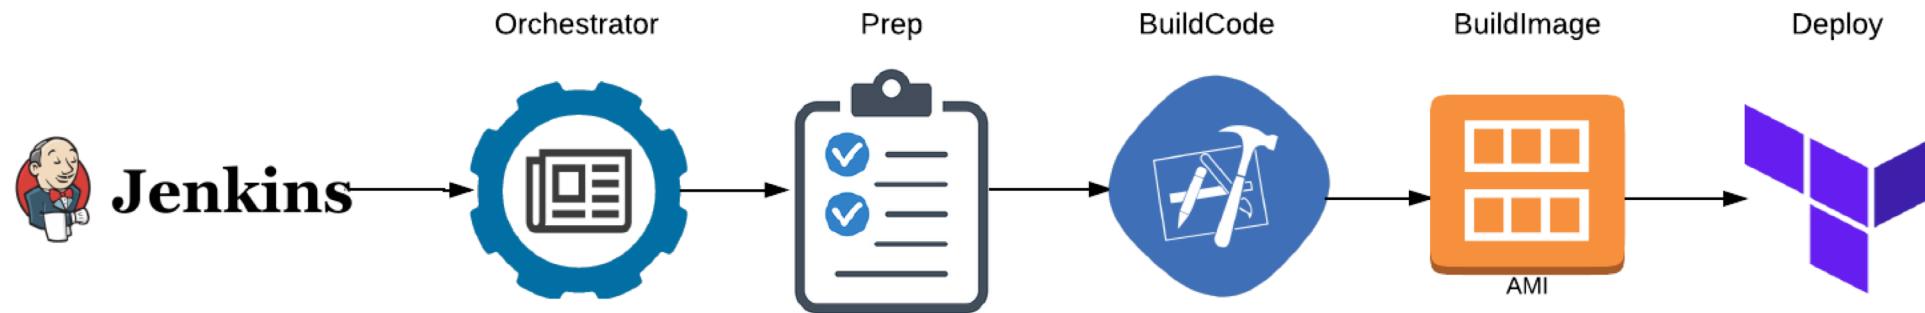


By The Numbers

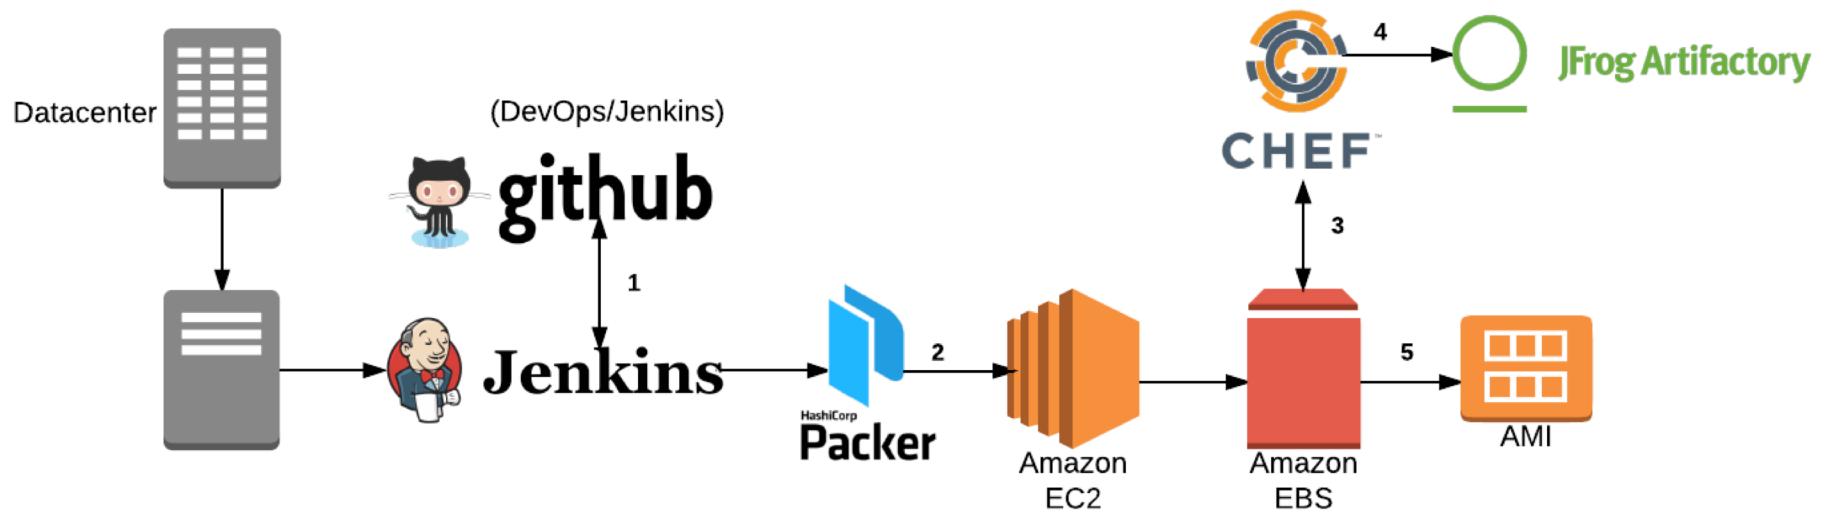
- 9 Websites
- 20+ Services
- 8 Databases
- 3 External Applications



The DevOps Pipeline



The DevOps Pipeline: BuildImage



The DevOps Pipeline: Packer

Packer is an open source tool for creating identical machine images for multiple platforms from a single configuration.

- Advantages
 - Lightweight, Highly performant
 - Runs on every major operating system
 - Creates images for multiple platforms in parallel (AWS, Vmware, Docker)



The DevOps Pipeline: Packer Tips & Tricks

I want to see more detail where the packer build is failing, what do I do?

- Debug
 - If you want to step through the build process, use the –debug flag in your packer command:

```
packer build -debug app.json
```

- Logging
 - You can also add logging to have the packer output written to a specified location:
 - Note: A log will be generated by setting PACKER_LOG equal to anything but “”

```
export PACKER_LOG_PATH="/path/to/packer.log"
```

```
export PACKER_LOG=10
```



The DevOps Pipeline: Packer Tips & Tricks

DEMO



The DevOps Pipeline: Chef

Chef – An open-source tool used for configuration management.

- Advantages

- Accelerating software delivery by allowing deployment of software within 1 hour of a commit.
- Integrated testing framework (Test Kitchen)
- Maintain a consistent configuration up the stack (Dev → QA → Staging → Prod)



The DevOps Pipeline: Chef

- Test Kitchen is a **testing** tool that can execute your CHEF recipes on one or more platforms in isolation. A driver plugin architecture is used which lets you run your recipes on various cloud providers and virtualization technologies.
- Setup
 - [Full Setup Tutorial](#)
- Types
 - Integration (*ServerSpec* based on *Rspec*)
 - Unit (*ChefSpec* based on *Rspec*)
 - Most common is **Integration**.

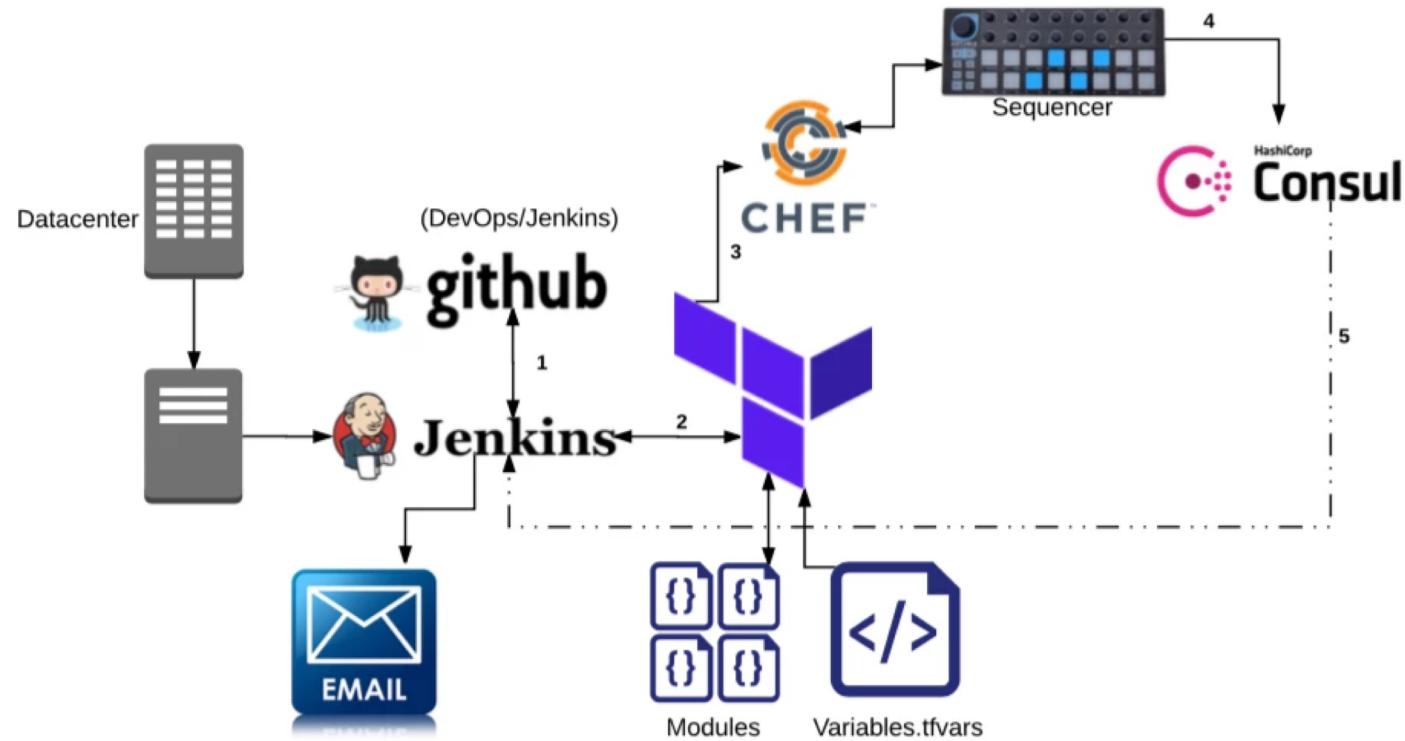


The DevOps Pipeline: Chef

DEMO



The DevOps Pipeline: Deploy



The DevOps Pipeline: Terraform

Terraform enables you to safely and predictably create, change, and improve infrastructure.

- Advantages
 - Cloud platform agnostic (AWS, Azure, GCP)
 - *Well...sort of....*
 - Code Readability (Written in GO, mostly API calls)



The DevOps Pipeline: Sequencer

What is Sequencer?

- Sequencer is essentially a script written in Ruby that takes values in and makes GET requests to Consul to check status on various components being deployed by Terraform.
- Sequencer is used to provide visibility into the deployment of Terraform components during provisioning (deployment).



The DevOps Pipeline: Sequencer

```
1 node['sequencer']['dependencies'].each do |dependency|
2   # Updates value to let user know we're waiting
3   http_request 'Set value to Waiting' do
4     message "WAITING on #{dependency}"
5     url "http://#{node['sequencer']['consul_address']}"
6     action :put
7   end
8
9   ruby_block 'Check dependencies' do
10    block do
11      require 'net/http'
12      require 'base64'
13      require 'json'
14      puts "#{dependency}"
15      http = Net::HTTP.new('consul.management.<redacted>.com', 80)
16      response = http.send_request('GET', "/v1/kv/sequencer/#{node['sequencer']['dependencies'][dependency]}")
17      puts response.body
18      puts response.code
19      json_response = JSON.parse(response.body)
20      # You'll notice the 2nd operator for the comparison (base64 encoded)
21      # Beware of this gotcha if you choose to refactor this
22      while json_response[0]["Value"] != Base64.encode64('FINISHED').strip
23        http = Net::HTTP.new('consul.management.<redacted>.com', 80)
24        response = http.send_request('GET', "/v1/kv/sequencer/#{node['sequencer']['dependencies'][dependency]}")
25        json_response = JSON.parse(response.body)
26        sleep 5
27      end
28    end
29  end
30 end
31
```

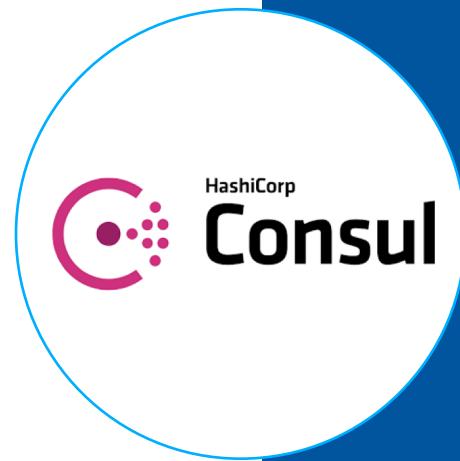


The DevOps Pipeline: Consul

Consul has multiple components, but as a whole, it is a tool for discovering and configuring services in your infrastructure. It provides several key features:
Service Discovery: Clients of Consul can provide a service, such as api or mysql , and other clients can use Consul to discover providers of a given service.

How do we use it?

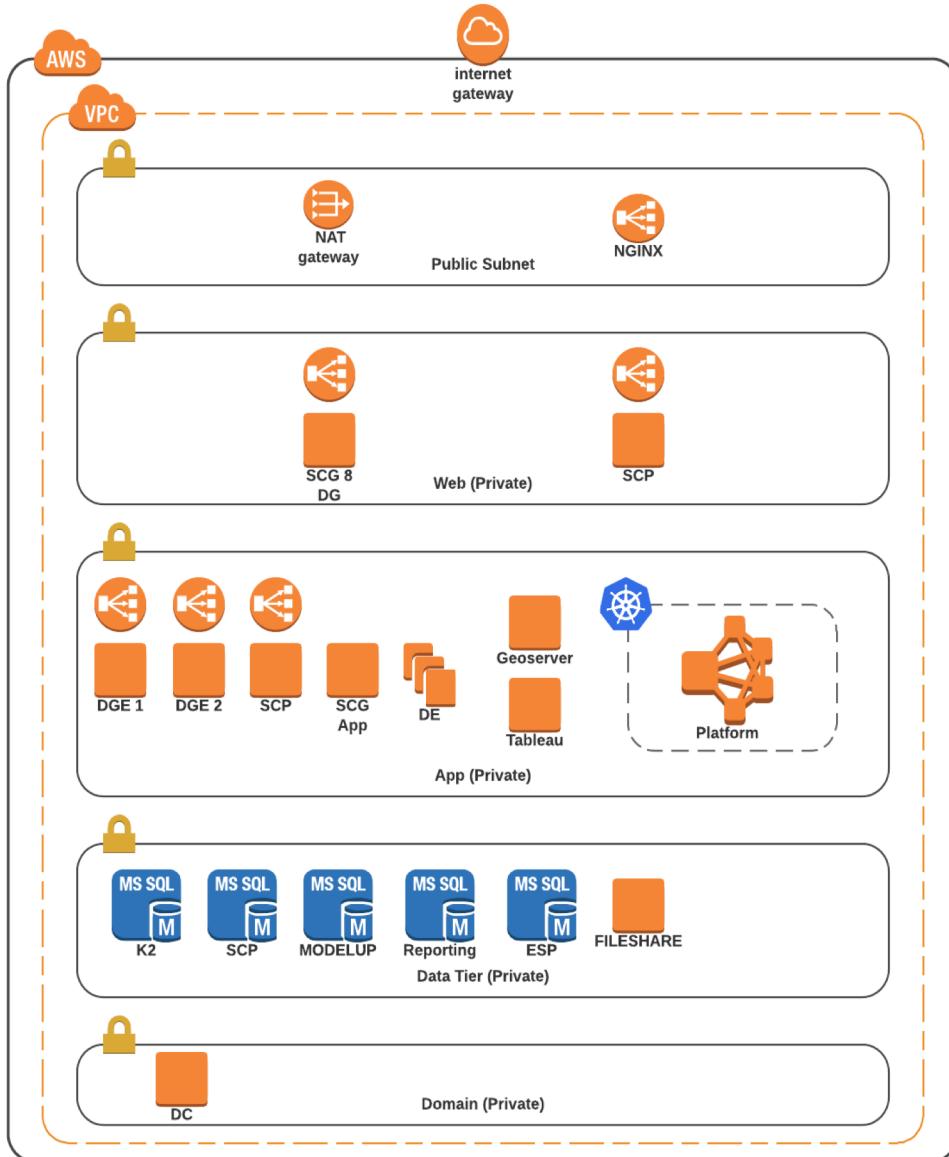
- Consul is simply a key/value store for us to store terraform state and be able to query whether a component deploy is running, waiting or deployed.



The DevOps Pipeline: Consultation



Basic Diagram



Single point of entry for all traffic; NAT gateway for necessary outbound requests

Separate tiers (Web, App, Data, Domain)

Load Balancers enable seamless scaling and increase security

Kubernetes provides scaling and recovery capabilities for platform services

Can be deployed in multiple AWS regions



The DevOps Pipeline: Jenkins

Jenkins is an open source automation server written in java. Jenkins helps to automate the non-human part of the software development process, with continuous integration and facilitating technical aspects of continuous delivery.

Declarative Pipeline Advantages:

- Provides richer syntactical features over Scripted Pipeline syntax.
- Designed to make writing and reading Pipeline code easier.



The DevOps Pipeline: Jenkins

```
1 // @Library('platform-helpers@feature/AddEFSProvisionertoClusterDeploy')
2
3 agentName = "linux"
4 agentLabel = "${-> println 'Right Now the Agent Name is ' + agentName; return agentName}"
5
6 pipeline {
7     ...
8     agent {label agentName}
9
10    parameters {
11        string(name: 'name', description: 'Select the Name for the cluster, you must use a fully qualified domain name (i.e. toolin')
12        choice(name: 'environment', description: 'The target AWS VPC to deploy the Kubernetes cluster into (required)', choices:'De')
13        string(name: 'infrastructure_branch', defaultValue: 'develop', description: 'branch name do EFS Provisioner deploy from')
14        choice(name: 'kubernetes_version', description: 'Select release of Kubernetes to deploy.', choices:'1.9.3\\n1.9.4')
15        string(name: 'node_count', description: 'Provide the number of nodes (1-100)', defaultValue: '2')
16        string(name: 'master_count', description: 'Provide the number of masters (1-4)', defaultValue: '1')
17        choice(name: 'dns_zone', description: 'Select the Hosted DNS Zone for the cluster. (supplychainsoft.com is Dev, supplychain')
18        choice(name: 'instance_size', description: 'Select VM size for the Instance size for Master and Node(s) (Dev/Prod).', choic
19        booleanParam(name: 'protect', description: 'Mark cluster as "Protected" to avoid deletion via Jenkins.', defaultValue: false
20    }
21
22
23    stages {
24        stage('Gathering information for Kubernetes cluster creation...') {
25            ...
26            steps {
27                dir('platform') {
28                    script {
29                        echo "Determining correct cluster configuration file and values from Environment parameter..."
30
31                        def kopsConfig = params.environment
32
33                        switch(kopsConfig){
34                            case 'Dev':
35                                echo "Found kops_config_dev.yaml, using as cluster config..."
36                                kopsConf= 'kops_config_dev.yaml'
37                                efssubnetIdAZ1='subnet-<redacted>'
38                                efssubnetIdAZ2='subnet-<redacted>'
```



Llamasoft®

Supply Chain By Design

Kubernetes

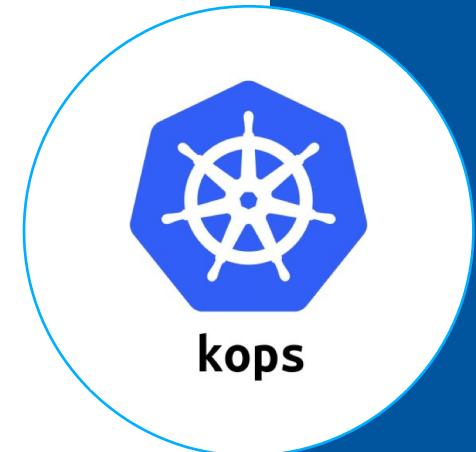


The DevOps Pipeline: kops

KOPS is an [official Kubernetes project](#) for managing production-grade Kubernetes clusters. Kops is currently the best tool to deploy Kubernetes to Amazon Web Services.

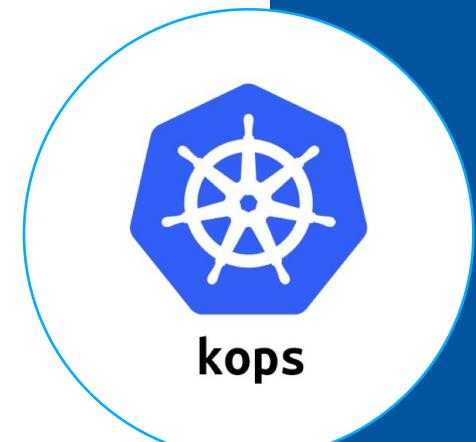
Key Features

- Rolling cluster updates
- Deploy clusters to existing virtual private clouds (VPC) or create a new VPC from scratch
- Provisions single or multiple master clusters



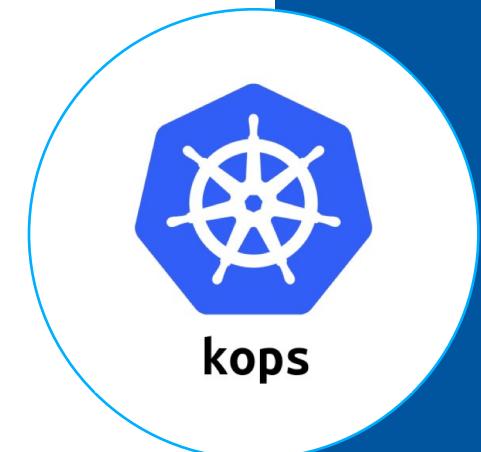
The DevOps Pipeline: kops

```
kops create cluster \
  --node-count 3 \
  --zones us-east-1a,us-east-1b \
  --dns-zone <redacted>.com. \
  --node-size m4.large \
  --master-size m4.large \
  --topology private \
  --vpc vpc-<redacted> \
  --network-cidr <redacted> \
  --networking kube-router \
  --cloud-labels "Environment=\"K8sDev\",Service=\"Kubernetes\",Type=\"INFRA\""
  --kubernetes-version 1.9.0 \
  --master-security-groups sg-<redacted> \
  --node-security-groups sg-<redacted> \
  dev.<redacted>.com
```



The DevOps Pipeline: kops

```
1 // @Library('platform-helpers@feature/AddEFSProvisionertoClusterDeploy')
2
3 agentName = "linux"
4 agentLabel = "${-> println 'Right Now the Agent Name is ' + agentName; return agentName}"
5
6 pipeline {
7     agent {label agentName}
8
9     parameters {
10         string(name: 'name', description: 'Select the Name for the cluster, you must use a fully qualified domain name (i.e. toolin')
11         choice(name: 'environment', description: 'The target AWS VPC to deploy the Kubernetes cluster into (required)', choices:'De')
12         string(name: 'infrastructure_branch', defaultValue: 'develop', description: 'branch name do EFS Provisioner deploy from')
13         choice(name: 'kubernetes_version', description: 'Select release of Kubernetes to deploy.', choices:'1.9.3\\n1.9.4')
14         string(name: 'node_count', description: 'Provide the number of nodes (1-100)', defaultValue: '2')
15         string(name: 'master_count', description: 'Provide the number of masters (1-4)', defaultValue: '1')
16         choice(name: 'dns_zone', description: 'Select the Hosted DNS Zone for the cluster. (supplychainsoft.com is Dev, supplychain')
17         choice(name: 'instance_size', description: 'Select VM size for the Instance size for Master and Node(s) (Dev/Prod).', choic
18         booleanParam(name: 'protect', description: 'Mark cluster as "Protected" to avoid deletion via Jenkins.', defaultValue: fals
19     }
20 }
21
22
23     stages {
24         stage('Gathering information for Kubernetes cluster creation...') {
25             steps {
26                 dir('platform') {
27                     script {
28                         echo "Determining correct cluster configuration file and values from Environment parameter..."
29
30                         def kopsConfig = params.environment
31
32                         switch(kopsConfig){
33                             case 'Dev':
34                                 echo "Found kops_config_dev.yaml, using as cluster config..."
35                                 kopsConf= 'kops_config_dev.yaml'
36                                 efssubnetIdAZ1='subnet-<redacted>'
37                                 efssubnetIdAZ2='subnet-<redacted>'
```

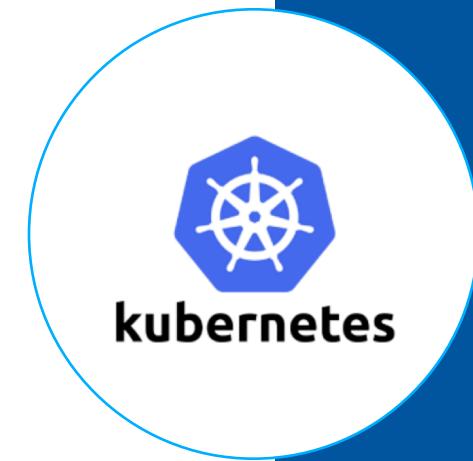


The DevOps Pipeline: Kubernetes

Kubernetes is a portable, extensible open-source platform for managing containerized workloads and services, that facilitates both declarative configuration and automation.

Advantages:

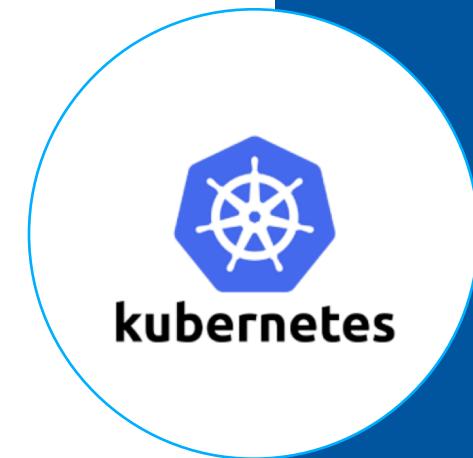
- High scalability
- Easier container management
- Makes rolling-updates *easy*



The DevOps Pipeline: Kubernetes

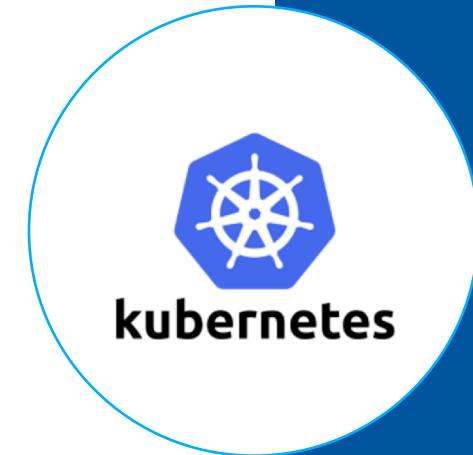
Helpful Hints for Deploying a Cluster:

- Use some sort of Pipeline mechanism for issuing the kops cli / kubectl commands.
- Make the Pipeline tool easy to use by creating a job (Maybe a Jenkins job) that passes parameters into the script you're using.
- Create a docker image with relevant tooling for your team to use. [Alpine Linux](#) is a nice lightweight option.



The DevOps Pipeline: Kubernetes

Codeshare



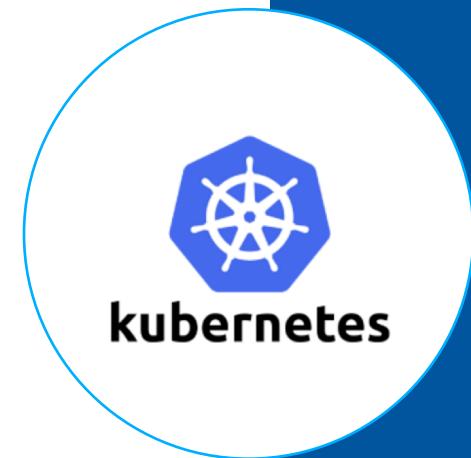
The DevOps Pipeline: Kubernetes

Learning:

- There is a lot of information out there about Kubernetes, some of it is old, some of it is incorrect, and some is from last nights release.
- Do yourself a favor and go take a [Pluralsight](#) course on it.
- Additionally, it's good to understand [Docker](#) prior to diving into Kubernetes, so maybe take a course on that as well.

Helpful Links:

- [Kubectl Cheat Sheet](#)
- [Kubernetes Tutorials](#)



The DevOps Pipeline: K8s Lessons Learned

Why isn't my microservice reachable?

Is the service running?

Issue: `kubectl get svc <service name> -n <namespace>`

If the service is running, can I see more detail?

Issue: `kubectl describe svc <service name> -n <namespace>`

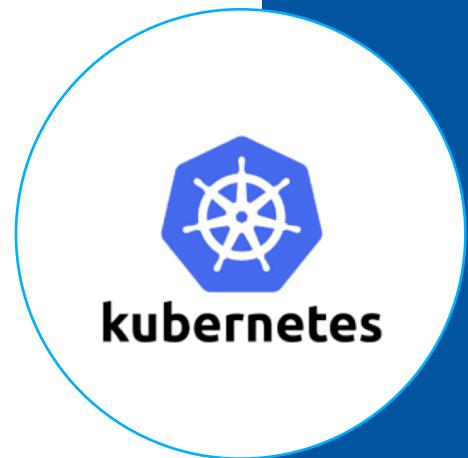
What about logs?

Issue: `kubectl get pods -n <namespace>` (Get pod name)

Issue: `kubectl logs pod <podname> -n <namespace>`

Can I tail the logs?

Issue: `kubectl logs -tail=<number of lines> <podname>`



The DevOps Pipeline: K8s Lessons Learned

According to the logs, my microservice is unable to make a connection to the database, what should I do to troubleshoot?

Are the environment variables (connection strings) correct?

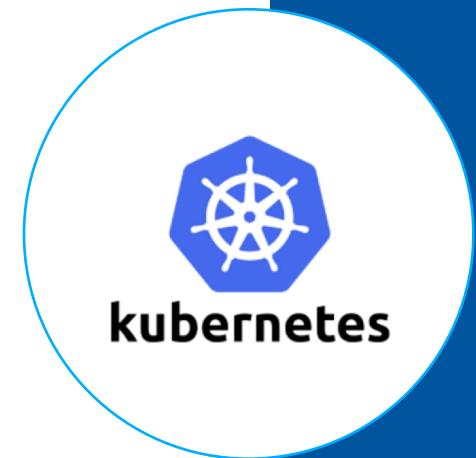
Issue: `kubectl get pods -n <namespace>` (Get pod name)

Issue: `kubectl describe pod <podname> -n <namespace>`

Can other microservices connect to the database, if so, can other microservices on the same Node as my microservice connect?

Issue: `kubectl get pods -n <namespace>` (Get pod name)

Issue: `kubectl exec -it <podname> bash/sh -n <namespace>`



Llamasoft®

Supply Chain By Design

What's Next?



The DevOps Pipeline: What's Next?

- Monitoring
 - Prometheus (Grafana, Alerts..etc)
 - Cloudwatch (Infra Alerts)
- Logging
 - Filebeat
 - Fluent-bit/Fluentd
- Tooling
 - Jenkins
 - Artifactory



The DevOps Pipeline: References

- [Chef Test Kitchen Tutorial](#)
- [Kubectl Cheat Sheet](#)
- Slack Communities
 - [Orchestecture](#)
 - [Chef](#)
 - [Kubernetes](#)



QUESTIONS?

Josh Lynn
email: josh.lynn@llamasoft.com
slack: @Lynny (Made in A2)

